Subset Sum Problem in optimization form: a parallel approach

Project for the GPU Computing course held by Professor Giuliano Grossi, University of Milan, February 2023

Alessio La Greca

Abstract

The *Subset Sum* with only positive values is a known problem in the field of computer science. The given input data is always:

- a vector A composed of n positive values. This represents a *multiset*, meaning that the same element can appear more than once;

- an objective value c.

In its basic form, the decision one, one is asked to tell whether there is a subset (which can as well be a sub-multiset) A' of A such that its elements, summed together, are equal to exactly c. This problem is known to be NP-complete. There is however an optimization form of the problem, where one is asked to find the value c' such that:

- c' is given by the sum of the elements of a certain sub-multiset A' of A;

- $c' \leq c$;

- There does not exist a value c'' given by a sub-multiset A'' of A such that $c'' < c \wedge c'' > c'$.

Basically, we want to find the sub-multiset whose values, summed together, get as close as possible to c. This problem is regarded as NP-hard. This particular problem, that will be of our complete interest throughout this report, and that we will simply call *SSO* (*SumsetSumOptimization*) from now on (as well as we'll call its decision form *SSD* (*SubsetSumDecision*)), comes in handy in different applications. For example, this problem can be used in a *Branch & Bound* algorithm for solving the 3D Knapsack Problem, as a procedure for computing the dual bound of a partial solution. There are different ways to tackle this problem: some give an exact result, while others use an heuristic approach. At the same time, different algorithms can be more or less efficient. Given that we'll focus only on **exact** algorithms, the goal of this report is to propose some procedures to solve the SSO problem by first studying sequential approaches, and then using parallel approaches to see what kind of speedup can be achieved. For the parallel algorithms, we'll leverage on a Nvidia GPU and its CUDA programming model.

Contents

1 Introduction

1.1 Disclaimer

Since the parallel algorithms are written in CUDA C, the reader shall be warned that some basic understanding of it is necessary to comprehend the implementation choices proposed in this report, altought the use of CUDA/Nvidia GPU specific features will be justified during the explaination of the relative algorithms. This means that an in-depth overview of things like the choice of the arrangement of the data in memory, the computations regarding thread indexes, and such, will be presented.

The code relative to the algorithms proposed here can be found in the same repository in which this report lies[1]. For further and deeper references, one can check the book[2] from which the CUDA C concepts used here have been learnt. We make the assumption that the reader is familiar with the basic CUDA terminology and concepts: each thread has a *threadIdx*, *blockIdx*, etc..

1.2 Terminology

From now on, we'll also refer informally to the vector of positive values, A, as *volumes*, while the target value c will be called *capacity*. This is in analogy to the sub-procedure used in a *Branch & Bound* approach for the 3D Knapsack Problem. The volumes vector A is considered to have lenght n, and it's a multiset, meaning that the same element can appear twice.

Given for example five given items, their volumes vector of length $n = 5$ would be:

$$A = \{2, 9, 9, 4, 7\}$$

Notice how the elements are not sorted nor in one way nor in the other. If not specified otherwise, we will assume that the order of the elements is not important, though *sometimes having them sorted is necessary to have a correct algorithm*. Chapter 5 will cover this in more detail.

Another representation of the volumes vector is possible: in such case, if we assume that m is the number of *different* values we have at our disposal, we can build a representation of the multiset based on two vectors:

- One is A, that this time contains only the different items. Considering the example shown before, we would have:

$$A = \{2, 9, 4, 7\}$$

- The other is B, a vector that counts how many times an element of A is present. In this case, it is:

$$B = \{1, 2, 1, 1\}$$

In this report, we will always stick with the first of the two representations, but in the future also algorithms based on the second one could be used.

1.3 Reference example

The main goal of this whole report is not to just show off some efficient code to solve this problem: also that, but it's also to study how one can approach a problem this hard through trial-and-error, refinement of different ideas and creativity, why not. In order to have a better grasp of what the algorithms are doing, a reference example can be useful. This is why we show here a toy instance that will be used to show the execution of the algorithm during their explanation. Note that sometimes we might need to show a more elaborated instance to better understand the algorithm, but unless differently stated, this one will be considered.

Our reference example has the following data:

- $A = \{3, 5, 8, 10\}$

- $c = 12$

This instance will be called *toy_4_12*, because $n = 4$ and $c = 12$. It's easy to see that this SSO instance has optimal value $11 = 3 + 8$.

1.4 Actual examples

To compare the performance of the algorithms, however, some practical examples need to be used. In this section we present the instances we will be dealing with in this report. The first instance will be called *custom_1*, according to the function *initialize_custom_1* (found in *myProject/common.h*) called by the *exec_main_0* command that can be found in the various makefile.

Volume value	Copies
1850000	2
120000	12
30000	10
1800	6
1250	4
400	2

- The total number of volumes is 36.

- The target capacity is 3690000.

- The best solution is 3606600.

Here, we listed the volumes stating also their number of copies, instead of enumerating all of them explicitly. Please note that we did this just to make the instance clearer to read: the actual representation used by the algorithms will be the first one discussed before, that is, an array of 36 elements. The first two entries will be 1850000, the next twelve will be 120000, and so on.

This instance should be thought as a meaningful one (that is, that could also be present in a real case scenario), although relatively simple. We will use this mainly for the sequential algorithms.

The second instance instead will be called *random_1*, name given by the fact that we ask to an auxiliary procedure to randomly produce n volumes within a given range, such that each of them is lower than the specified capacity. This instance, that can be launched from basically every makefile present in the repository, is the *exec_main_1*, where the arguments (in their order of appearance) mean:

1. 0: specifies that we want to use the auxiliary procedure for generating random volumes.

2. 1: once again specifies that, in the auxiliary procedure, the volumes should be really randomly generated, and not taken from a pre-determined *for* loop that would otherwise generate them.

3. 100: the number of volumes we want for this instance.

4. 1000000000: the capacity (one billion).

5. 1: the seed for the random generation. This will ensure us that, although the procedure randomly generates the volume values, they will always be the same for every execution.

This instance, differently from the previous one, has no real meaning, and is such that actually the optimal solution requires to take all of the volumes. But, being it very large both in termis of n and c, it will be useful to compare fast GPU-based implementations.

For each algorithm, we will execute each considered instance **five** times (if not specified otherwise), reporting the minimum, maximum and average time of execution. For the GPU, other profiling data could be presented.

2 CPU (Sequential) algorithms: exhaustive search

2.1 SSO: Exhaustive search

The ol'good exhaustive approach, although obvious, is the first that must be presented, both to have a worst-case performance algorithm to which we'll compare the following ones, and to have a procedure that no matter what will give us the optimal value, which can then be used to make sure that the subsequent implementations are correct.

The exhaustive search can be easily performed with a recursive algorithm. The algorithm produces all the 2^n strings of bit possible, where n is the number of volumes. In a particular string of n bits, let's call it S, we have $S[i] = 1$ if the i-th element of the vector A of volumes is considered in the current tentative solution, 0 otherwise. The following table shows the tentative solutions produced for the *toy_4_12* instance:

Bit string	Value	Feasible?	Optimal?
0000	$0 + 0 + 0 + 0 = 0$	yes	no
0001	$0 + 0 + 0 + 10 = 10$	yes	no
0010	$0 + 0 + 8 + 0 = 8$	yes	no
0011	$0 + 0 + 8 + 10 = 18$	no	no
0100	$0 + 5 + 0 + 0 = 5$	yes	no
0101	$0 + 5 + 0 + 10 = 15$	no	no
0110	$0 + 5 + 8 + 0 = 13$	no	no
0111	$0 + 5 + 8 + 10 = 23$	no	no
1000	$3 + 0 + 0 + 0 = 3$	yes	no
1001	$3 + 0 + 0 + 10 = 13$	no	no
1010	$3 + 0 + 8 + 0 = 11$	yes	yes
1011	$3 + 0 + 8 + 10 = 21$	no	no
1100	$3 + 5 + 0 + 0 = 8$	yes	no
1101	$3 + 5 + 0 + 10 = 18$	no	no
1110	$3 + 5 + 8 + 0 = 16$	no	no
1111	$3 + 5 + 8 + 10 = 26$	no	no

The time complexity of this algorithm is $\Theta(2^n)$, since all the possible solutions are visited. For the same reason, the space complexity of a recursive approach is $\mathcal{O}(n)$.

2.2 SSO: Smarter exaustive search

The exhaustive approach starts with an *accum* variable, initialized to 0, that keeps track of the current sum produced do far. At each call, two recursive calls are made: in one, we add to the solution the next volume in A. In the other, we skip it. We can avoid visiting

branches, in the recursion tree, that are made only of unfeasible solutions by saying, at a certain call:

- If adding the current volume to the current partial solution does not exceed c, then both calls must be performed.

- But if it does, simply skip it: every solution that sums the current partial solution with this volume will be unfeasible.

This basically allows us to perform the exhaustive search only in the feasible region of the search space, since all the (partial) unfeasible solutions are immediately discarded. The following table shows the tentative solutions produced in this case:

Bit string	Value	Optimal?
0000	$0 + 0 + 0 + 0 = 0$	no
0001	$0 + 0 + 0 + 10 = 10$	no
0010	$0 + 0 + 8 + 0 = 8$	no
0100	$0 + 5 + 0 + 0 = 5$	no
1000	$3 + 0 + 0 + 0 = 3$	no
1010	$3 + 0 + 8 + 0 = 11$	yes
1100	$3 + 5 + 0 + 0 = 8$	no

The time complexity of this algorithm is $\mathcal{O}(2^n)$, since only in the worst case all the possible solutions are visited. The space complexity is $\mathcal{O}(n)$ once again.
The code for this algorithm can be found in $myProject/DIR\_\_\_CPU\_exhaustive\_opt$.
Its performances are reported in the following table.

Instance	Min	Max	Avg
custom_1	89.788956 s	91.011294 s	90.330596 s

Actually, in an attempt to make the algorithm more efficient, we implemented it in a way that combines two approaches: a constructive and a destructive one. This came from an intuition: sometimes, instead of starting from the value 0 and "enlarging" it as close as possible to the target capacity, it can be better to start from the sum of all the volumes and removing some of them, until we reach a feasible solution.
The algorithm proceeds as follows:

1. Given the volumes, their number and the capacity, the mean volume is computed (the sum of all the volumes divided by their number).

2. The capacity divided by this mean gives us an approximation of the number of volumes needed to fill the capacity. If this value is greater than half of the total number of volumes, the destructive approach is used. Otherwise, the constructive one is.

This is an approach that is good in theory but can be arbitrarily bad in practice. Much depends on the distribution of the volumes and the capacity value. Still, it was necessary pointing out this implementation detail, that could be further addressed in the future.

```
SUBSET-SUM-OPTIMIZATION(A, b)
      for i ← floor(log₂b) downto 0 do
            A ← A + { 2ⁱ }
      for i ← floor(log₂b) downto 0 do
            A ← A - { 2ⁱ }
            if not SUBSET-SUM-DECISION(A, b) then
                  b ← b - 2ⁱ
      return b
```

Figure 1: Chad Parry's Algorithm for the SubsetSumOptimization problem

3 CPU (Sequential) algorithms: Chad Parry's algorithm

A completely different approach for solving the SSO problem comes from an algorithm proposed by Chad Parry[3], where $b = c$:

The idea is first of all to enlarge the set of volumes with fictitious volumes that are all powers of 2, and then to solve multiple times the SSD problem in order to determine the optimal solution for the original SSO problem. This effectively reduces the SSO problem to the SSD problem, altought this must be solved multiple times. This means that if we can find an efficient way to solve the SSD problem, we automatically have an efficient way to solve the SSO problem. That said, we will now focus on different ways in which we can solve the SSD problem.

3.1 Important premise

Our main concern here is **efficiency**, or better **an efficient way to solve the SSO problem**. Since Chad Parry's algorithm for the SSO problem relies on an iterated application of a procedure that solves the SSD problem, the overall complexity of the following algorithms is given by the complexity of the procedure that solves the SSD problem multiplied by the number of times that procedure is executed ($\lfloor log_2(c) \rfloor$).

...and even so the computations wouldn't actually be precise, since each time the SSD procedure is executed, the input data are smaller than the previous execution (there is one less item): in short, to actually make an exact computational complexity study, we would need to consider all these things. But is it really necessary? We don't think so, since the only thing that we can manipulate in the Chad Parry's algorithm is the SSD procedure. We can easily see that, the more efficient this specific function is, the more the overall algorithm will be. So, we will focus on the time and space complexity of the proposed SSD procedure.

3.2 SSD: Exhaustive search

Nothing stops us from taking the exhaustive approach used for the SSO problem in Section 2.1 and use it to solve the decision form of the problem. The only difference is that this time we begin with the intention to fully explore the search space, by considering every possible string of n bits, but, thanks to the short-circuit logic of the *or* operator of the C language, we stop our search as soon as a solution equal to c is found (otherwise, we fully explore the search space).

The time complexity of this algorithm is $\mathcal{O}(2^n)$, while the space complexity of a recursive approach is $\mathcal{O}(n)$.

The code for this algorithm can be found in *myProject/DIR___CPU_cp_exhaustive_dec*.

In an attempt to run this algorithm to collect its performance on the *custom_1* instance, the execution time exceeded 3 full hours, without coming to an end. The code is correct, and terminates giving the expected output for smaller instances: this one in particular, however, is way too big for it to be handled. This was to be expected since we basically repeat an exponential search multiple times. Let's than see how to address this problem.

3.3 SSD: Dynamic Programming

An exhaustive search is intrinsically inefficient: although such an algorithm can be used when the number of volumes is pretty low, as soon as this number rises a bit, the complexity becomes worse and worse. So, what we propose here is a Dynamic Programming approach to solve the SSD problem. Altough we could just present the most efficient algorithm we've found, we will instead show all the approaches we've come up with, for reasons stated in Section 1.3.

But before that, we need to state the general idea: a Dynamic Programming way to solve the SSD problem is to build a binary matrix DP with $n + 1$ rows and $c + 1$ columns, and to progressively fill each row, one after the other. Each cell of each row can be fully determined by at most two cells of the previous row.

At row i and column j, the cell DP[i][j] is *true* if there exist a subset of the first i items in A that can sum up to j, *false* otherwise. Note that the order of the elements in A does not affect the final result.

- The first row covers the case in which no item is considered. This leads to an initial row that has the first element equal to *true*, since summing no items gives 0, that fits exactly a capacity of 0, and *false* for every other cell of the row, because without volumes we can't fill a positive capacity.

- Given any of the subsequent rows, we distinguish two cases:

  1. $A[i] > j$, that is, the item $A[i]$ associated to the current row i is greater than the current capacity j, so it can't possibly fit. In this case, we simply copy the previous case, since we can't do much else: $A[i] = A[i-1]$.

  2. $A[i] \leq j$. This time, item $A[i]$ fits in the capacity j, but this *does not mean that it is part of the solution*. In fact, whether capacity j was already reached without the support of this item (DP[i-1][j] = *true*), adding it to the input data does not change the result. However, if instead DP[i-1][j] = *false*, then we can try to see if the sub-problem of the row before that had capacity $j - A[i]$ admitted a postive answer, because if it did, then this does too. In short:
  $DP[i][j] = DP[i-1][j] \vee DP[i-1][j - A[i]]$

By taking the element DP[n][c] (last row, last column), we can find out whether there is a positive answer to the original SSD problem or not.

3.3.1 DP: Version 1

The first implementation of this approach is extremely straightforward: just apply the algorithm as we have described it. We initialize the first row, and then compute all the following ones based on the one that precedes it. To do so, we need to allocate $n + 1$ rows of length $c + 1$ each. Since our code is developed in CUDA C, we could declare a row as an integer array. However, the type *int* in C occupies 32 bits, which is a waste for just a binary information. This, together with the fact that there isn't a built-in *boolean* type in C, led us to the choice of representing a cell of the matrix as an *unsigned char*: it will be equal to 1 (*true*) when the first i objects admit a subset that sums up to j, 0 (*false*) otherwise. This is because the unsigned char built-in type occupies just one byte of memory, the least amount of information a standard C variable can hold. The following table is the resulting matrix produced by this algorithm for the *toy_4_12* instance.

Actually, in order to be coherent with Chad Parry's algorithm, we should build multiple matrixes: one in which we consider the volumes given by the problem plus the fictitious volumes $\{8, 4, 2, 1\}$, then another with the same volumes except 8, and so on. However, to simplify the presentation of the concepts, since we just want to compare different SSD algorithms from an efficiency point of view, we'll just apply the DP algorithm to the decision version of the original *toy_4_12* reference example.

	0	1	2	3	4	5	6	7	8	9	10	11	12
0	T	F	F	F	F	F	F	F	F	F	F	F	F
3	T	F	F	T	F	F	F	F	F	F	F	F	F
5	T	F	F	T	F	T	F	F	T	F	F	F	F
8	T	F	F	T	F	T	F	F	T	F	F	T	F
10	T	F	F	T	F	T	F	F	T	F	T	T	F

Of course, the column index dictates the target capacity c of the considered sub-problem, while the number at the beginning of each row represents the value of the new volume considered.

The time complexity of this algorithm is $\Theta(n \times c)$ (to be precise it would be exactly $(n + 1) \times (c + 1)$), and the space complexity is exactly the same, since we allocated $n + 1$ rows of $c + 1$ elements each. Note that in this case the space complexity coincides exactly with the number of bytes used, since each cell is represented, in our CUDA C implementation, by an unsigned char.

The code for this algorithm can be found in *myProject/DIR___CPU_cp_dp*, in the file *CPU_dp.cu*.

Its performances are reported in the following table.

Instance	Min	Max	Avg
custom_1	10.637604 s	10.928177 s	10.741072 s

3.3.2 DP: Version 2

For a problem like this, it's not that difficult to have instances with a very high capacity. Even billions are totally common. If we consider that this high capacity value gets multiplied by the number of items, things get worse, both time-wise but most importantly

space-wise. We say "most importantly" because, even though the time complexity is our main concern, if we don't have enough memory to solve the problem, we can't do much. And $\Theta(n \times c)$ bytes seem a lot, can't we do something better? Of course we can.

Since every line other than the first is completely determined by the preceeding one, we can allocate just two rows:

1. Initially, fill the first one as we know: *true* in the first cell, *false* in every other one.

2. Set a variable $i = 0$, that represents the current number of items we are considering.

3. Repeat the following steps until $i = n$ (every volume has been considered):

   (a) $i = i + 1$;
   (b) Compute the second row from the first one, by considering only the first i items;
   (c) Overwrite the first row with the second.

The following tables represent the subsequent states of the matrix throughout the execution of the algorithm on the *toy_4_12* instance:

	0	1	2	3	4	5	6	7	8	9	10	11	12
0	T	F	F	F	F	F	F	F	F	F	F	F	F
3													

	0	1	2	3	4	5	6	7	8	9	10	11	12
0	T	F	F	F	F	F	F	F	F	F	F	F	F
3	T	F	F	T	F	F	F	F	F	F	F	F	F

	0	1	2	3	4	5	6	7	8	9	10	11	12
3	T	F	F	T	F	F	F	F	F	F	F	F	F
5	T	F	F	T	F	T	F	F	T	F	F	F	F

	0	1	2	3	4	5	6	7	8	9	10	11	12
5	T	F	F	T	F	T	F	F	T	F	F	F	F
8	T	F	F	T	F	T	F	F	T	F	F	T	F

	0	1	2	3	4	5	6	7	8	9	10	11	12
8	T	F	F	T	F	T	F	F	T	F	F	T	F
10	T	F	F	T	F	T	F	F	T	F	T	T	F

The time complexity is unchanged, we still need $\Theta(n \times c)$ steps to compute the final result, but the memory now depends only on the value of c, and it's always $\Theta(2 \times c)$. One more thing to take into account, however, is the overhead caused by copying the second row into the first one at each step: from a practical point of view, it influences the overall time complexity of the procedure, making it worse then the previous version.

The code for this algorithm can be found in *myProject/DIR___CPU_cp_dp*, in the file *CPU_dp.cu*.

Its performances are reported in the following table.

Instance	Min	Max	Avg
custom_1	15.786004 s	16.290045 s	15.973391 s

3.3.3   DP: Version 3

With the last paragraph, we've seen an implementation that improves memory usage, but worsens time complexity because of the overhead brought by the copy operations. However, it allowed us to reason about one thing: to compute a row, only the previous one is necessary. We will now take this argument one step forward, improving both time and space complexity.

It is true that row i can be fully determined by row $i-1$. But what actually happens is that the j-th element of row i depends only on two elements of row $i-1$: the j-th one and the $(j-A[i])$-th one. Or at least, this is true for every $j \geq A[i]$, at row i. This could suggest us that copying the first $A[i]$ elements of row $i-1$ into row i is redundant.

At the same time, up until now we've filled row i starting from the leftmost element to the rightmost one. Nothing stops us from doing it also the other way round. But why doing so? Because in this way we no longer need two data structures to store the old and new row, and we can do everything in-place. This does not affect correctness. Each element is in fact dependent only on elements on its left. Given this, modifying in-place the contents of the array starting from the last element makes the array such that, at a certain step i, it can be divided into two parts:

- A left one, composed of elements that consider only the first $i-1$ volumes. For each of these elements, it's true that the only entries necessary to compute the new value (the one that considers also the new i-th volume) are the element itself and the one $A[i]$ positions on its left.

- A right one, with elements that instead consider the first i volumes. These values are already updated, and won't be used by any of the left part, since they only care about elements on their left.

So, we have found a procedure that needs always one row (an array of $c+1$ entries), and that does not have the overhead of copying elements, since every value update is done in-place. At the end, the result can be red in the last position fo the array. Given the *toy_4_12* example, the following are the states of the array data structure of length $c+1$ during the execution of this algorithm. The red elements represent the values that have actually been computed, while the black one were not touched:

	0	1	2	3	4	5	6	7	8	9	10	11	12
0	T	F	F	F	F	F	F	F	F	F	F	F	F

	0	1	2	3	4	5	6	7	8	9	10	11	12
3	T	F	F	T	F	F	F	F	F	F	F	F	F

	0	1	2	3	4	5	6	7	8	9	10	11	12
5	T	F	F	T	F	T	F	F	T	F	F	F	F

	0	1	2	3	4	5	6	7	8	9	10	11	12
8	T	F	F	T	F	T	F	F	T	F	F	T	F

	0	1	2	3	4	5	6	7	8	9	10	11	12
10	T	F	F	T	F	T	F	F	T	F	T	T	F

The time complexity is $\mathcal{O}(n \times c)$, because this time we don't do exactly c operations at each step, but a bit less, depending on the current volume. The space complexity is $\Theta(c)$, possibly the smallest it can be.

The code for this algorithm can be found in *myProject/DIR___CPU_cp_dp*, in the file *CPU_dp.cu*.

Its performances are reported in the following table.

Instance	Min	Max	Avg
custom_1	7.230589 s	7.449816	7.347593 s

Of course, for instances with a very high c, we might still have memory problems, as we need to allocate c bytes of memory for this approach, but this is the best we can do with a Dynamic Programming sequential approach. The next step will be to translate this idea in a parallel algorithm.

4 GPU (Parallel) algorithms: Chad Parry's algorithm

Chad Parry's algorithm for the SSO problem, as we've seen in the previous chapter, is composed of multiple applications of a SSD procedure, that is up to us to implement. The most efficient implementation (time-wise at least) that we have seen is the one based on Dynamic Programming. This particular implementation is interesting also because it consists of many independent and simple operations. While it's true that each row needs the previous one to compute its entries (so, the rows can't be computed in parallel, but sequentially), the entries themselves are completely independent one another, and just need at most two elements of the previous row. So, they can be computed in parallel.

In a very abstract way, what each entry of row i does to compute its value is:

    if $j < A[i]$ **then**
        $DP[i][j] = DP[i-1][j]$
    else
        $DP[i][j] = DP[i-1][j] \vee DP[i-1][j - A[i]]$
    end if

This is a very simple procedure, and as such we can think of a CUDA application in which we run $c+1$ threads, where each of them deals with a particular entry of the DP array. The entry it deals with can be computed using its *threadIdx* and *blockIdx*. Both for simplicity and similarity, we can think to organize the threads in 1-dimensional blocks. This allows us to compute the thread index inside the grid (and so the element of the DP array it deals with) as $idx = blockIdx.x * blockDim.x + threadIdx.x$. That being said, we can no longer resort on using just one array of $c+1$ entries as we did in the 3rd version of the DP sequential implementation. This is because every element of row i must read at most two elements of the previous row, that must be kept consistent throughout the computations of this new row. So, the space complexity of the following algorithms for solving the SSD with a Dynamic Programming approach will have a space complexity of $\Theta(2 \times c)$. Their differences will lie in the time complexity.

The general scheme of the following algorithms is basically always the same:

1. Allocate, on the GPU, two arrays of length $c+1$: let's call them *row_0* and *row_1*. **Once again, to avoid wasting memory, these will be *unsigned char* arrays. This is something to keep in mind, expecially when we'll talk about shared memory.**

2. Launch a first kernel, *kernel_a*, to initialize *row_0* as expected: *true* as first element, *false* for every other one.

3. For c times, launch a second kernel, *kernel_b*, that computes the i-th row from the $i-1$-th one.

4. Finally, retrieve the result from the last entry of the last row computed.

Since the complexity of the initialization kernel *kernel_a* is insignificant compared to the iterated application of *kernel_b*, we will focus only on the implementation details of this last one (also because all the *kernel_a* implementations are actually more or less the same). Moreover, if up until now we've focused on the asymptotic notations of $\mathcal{O}$ and Θ to describe the performance of the algorithms, we will now instead examine their efficiency in CUDA terms.

Once again, we won't give immediately the best algorithm achieved: we will examine different ideas, by evolving each of them when possible in order to understand how one can exploit all of the Nvidia GPU features to achieve better and better performances.

4.1 SSD: Basic v1

The first approach that we propose is very naive and inefficient, but its simplicity will allow us to correctly reason about how to improve it.

We first allocate a total of three unsigned char arrays of length $c + 1$:

- The first is on the host.

- The second and the third, respectively *row_0* and *row_1*, are on the device.

After having initialized the host array with the starting values, we compute each following row as follows:

1. We copy the content of the host array into *row_0*;

2. We run *kernel_b* that computes the new row, *row_1*, from *row_0*;

3. We copy the content of *row_1* into the host array, and repeat from point 1.

In order to be a bit more efficient, at each iteration we check if the last element of the host array is *true*: if it is, we don't need to compute the following rows and we can immediately return a positive outcome.

The main drawback of this approach is that it makes a lot of copies from the host to the device and viceversa, worsening the overall performance. It's no secret that in CUDA programming the host-to-device and device-to-host data copy is an important bottleneck for many applications.

The code for this algorithm can be found in *myProject/DIR___GPU_cp_dp*, in the file *GPU_dp_basic.cu*.

Its performances are reported in the following table. Since in CUDA an important factor is also the block dimension, from now on we will execute the algorithm with multiple block configurations and see how the performances change. The ones we chose are 128, 256, 512 and 1024, since are usually the most meaningful.

Instance	Block dim.	Min	Max	Avg
custom_1	128	1.710427 s	1.843904 s	1.760543 s
custom_1	256	1.696415 s	1.843320 s	1.761409 s
custom_1	512	1.808097 s	1.891201 s	1.843313 s
custom_1	1024	1.721849 s	1.870658 s	1.789238 s
random_1	128	922.991272 s	922.991272 s	922.991272 s
random_1	256	932.092773 s	932.092773 s	932.092773 s
random_1	512	925.151550 s	925.151550 s	925.151550 s
random_1	1024	925.771912 s	925.771912 s	925.771912 s

Please note that the *random_1* instance was run only once per block configuration. Given how poor its time is, it wouldn't make much sense to average five different runs to make sure that this apporoach is quite terrible.

Despite the massive number of copies from host to device and viceversa, this kernel still performs very well with the *custom_1* instance. But this is no surprise: we already stated how this is actually a very simple one. In fact, the limited capacity is such that not much data needs to be transfered, and not for many times either, since $n = 36$, which is a pretty low value. The same, however, cannot be said for the *random_1* instance, that instead requires to transfer many times (100) **a lot** of bytes (almost 1 GB, given the capacity of one billion). Moreover, please note that we purposefully made this instance such that the optimal value is strictly lower than the capacity. In this way, the algorithm won't stop midway during the execution returning a positive answer, but instead will have to compute exactly n rows, which is the worst case scenario. From now on, during this chapter, we will focus on the *random_1* instance, since the *custom_1* would give very little detail about the actual improvement of the following algorithms, given that its execution time is already so low. The random instance, instead, is solved *so badly* from this implementation that it will be a good reference for the other proposed approaches.

4.2 SSD: Basic v2

The very first improvement we can come up with is one that doesn't perform so many memory operations from host to device and viceversa. After all, we don't really care bringing the intermediate rows to the host (except maybe for the last element that could be *true*), we just care about (the last element of) the last one. So, what can be done?

We could think of performing device to device memory copy operations, but this is rather similar to the approach seen in Section 3.3.2, which wasn't very efficient. To avoid unnecessary copy operations, we can just swap the roles of the two device arrays, *row_0* and *row_1*, when a new row needs to be computed. The overall algorithm would be as follows:

1. Allocate on the device two arrays of $c + 1$ unsigned char: our old and good *row_0* and *row_1*;

2. Initialize *row_0* with *kernel_a*;

3. label *row_0* as the "*old row*", and *row_1* as the "*new row*";

4. For n times, do:

   (a) Compute the *new row* from the *old row* using *kernel_b*;

   (b) Swap the *old row* and *new row* labels, so that the new one becomes the old one (for the next iteration) and viceversa;

14

(c) Check if the last element of the row just computed is *true*, and if it is, stop the loop and return *true*.

In this way we never perform any copy at the cost of simply remembering which row is the old one and which is the new one.

The code for this algorithm can be found in *myProject/DIR___GPU_cp_dp/DIR___GPU_cp_dp_basic*, in the file *GPU_dp_basic.cu*.

Its performances are reported in the following table.

Instance	Block dimension	Min	Max	Avg
custom_1	128	boh	bih	bah
custom_1	256	boh	bih	bah
custom_1	512	boh	bih	bah
custom_1	1024	boh	bih	bah
random_1	128	boh	bih	bah
random_1	256	boh	bih	bah
random_1	512	boh	bih	bah
random_1	1024	boh	bih	bah

Also forse altre considerazioni sui trasferimenti di memoria?

Please also note that the following sections of this chapter will start from this implementation, trying to improve it with different CUDA features.

4.3 SSD: Unrolling

The *unrolling* is a common technique in CUDA programming, that aims to improve the performance of kernels. The main idea is to take a for loop performed in a thread, and reduce the number of times it is executed by writing its body multiple times. This has two major advantages:

- We have a reduced instruction overhead, because the maintenance operations ($i < x; i++$) are performed less times.

- If the unrolled body now contains independent memory load/store operations, these can be scheduled simultaneously by the hardware, as memory latency can be better hidden.

That being said, the unrolling doesn't need to have a loop to start with. The idea is generally implemented by saying: take a thread and make it perform multiple independent operations. In our case, this is absolutely doable: each thread can handle the computation of k values for row i, instead of just one. This reduces by a factor k the overall number of threads needed, and increments the number of operations each thread must perform. However, being these memory independent operations, our hope is to achieve a speedup thanks to the fact that these will be scheduled efficiently and simultaneously by the scheduler.

The only (human) problem are the indexes: we must sort out, for each thread, the k indexes of the elements of row i that must be computed, but in such a way that each thread handles k elements that are different from those associated to every other thread. Now, given this constraint, generally speaking these indexes can be chosen arbitrarily. However, in order to be efficient, the accesses made to the global memory (where *row_0* and *row_1* lie) should be **aligned** and **coalesced**.

In the CUDA programming model, 32 threads with consecutive indexes work in a *SIMT* fashion to achieve parallelism, and they are collectively called *warp*. This affects the way instructions and memory operations are performed. In particular, when the threads of a warp want to access global memory, they group together their *individual requested addresses* into one *memory access request*. This memory access request is serviced in one or more **memory transactions**, that can be seen as the "currency" for any global memory access request. The less memory transactions a memory access request needs to be satisfied, the more efficient the access becomes.

Now, a memory transaction can be either of 128-bytes or 32-bytes, depending on whether the L1 on-chip cache is used or not. The important thing to note is that **these bytes are always contiguous** and **their initial address is always a multiple of 128 or 32**. To service even a memory access to just one byte, at least 32 consecutive bytes need to be read. This is why a memory access request should be composed of individual requested addresses that are *aligned* and *coalesced*, because in this way the number of necessary memory transactions is reduced.

- *Aligned* means that the first individual address requested by a warp is an even multiple of 32 or 128, respectively the granularity of the L2 device cache or L1 on-chip cache.

- *Coalesced* means that all the addresses requested by the threads of the warp are contiguous.

When both these assumptions hold for a memory request, we can be sure that the access will be efficient. The most common case is the one in which the threads of a warp read one integer each (4 bytes), and these bytes are such that:

- The first address is an even multiple of 32/128;

- All of them are adjacent in memory (for example, they are consecutive elements of an integer array).

Depending on the particular application, enabling the L1 cache can provide better performance or not. However, talking about our specific problem: we have had threads that don't read 4 bytes each, but one byte each. We have to take this into account when reasoning about how to arrange the memory indexes to these threads when performing the unrolling.

4.3.1 Unrolling factor 8

Let's then try to implement the unrolling approach to our kernels. To do that, we first have to choose an *unrolling factor*, that in this case translates to the value of k, that is how many elements of the new row a thread should compute. We first propose the value $k = 8$, since it's a pretty standard one.

Once again, after *row_0* has been initialized, we can modify the structure of *kernel_b* in such a way that each thread deals with 8 consecutive elements of the new row.

Each thread computes its first index as:

$$idx = blockIdx.x * blockDim.x * 8 + threadIdx.x * 8$$

and obtains the seven following ones by summing 1, 2, ..., 7 do *idx*. Each of these threads reads the corresponding value of the previous row and if necessary the one that is $A[i]$ positions before that. In this way, the 32 threads of a warp issue a collective memory request of 256 or 512 bytes to the global memory, such that:

- The index of the first address is always a multiple of 256, that is both a multiple of 32 and 128, so the access is aligned.

- All these bytes are contiguous, which means that the access is also coalesced.

The code for this algorithm can be found in
myProject/DIR___GPU_cp_dp/DIR___GPU_cp_dp_uroll, in the file *GPU_dp_uroll.cu*.
Its performances are reported in the following table.

Instance	Block dimension	Min	Max	Avg
custom_1	128	boh	bih	bah
custom_1	256	boh	bih	bah
custom_1	512	boh	bih	bah
custom_1	1024	boh	bih	bah

In this case, the use of L1 cache can improve a bit the overall performances. To see that, run first the application with the L1 cache enabled (using *-Xptxas -dlcm=ca*), then run again with it disabled (using *-Xptxas -dlcm=cg*).
Once again performance bla bla bla e oh, giusto metti a confronto diverse dimensioni di blocco

4.3.2 Unrolling factor 16

The improved performances of the last kernel suggest us to try to push even further the unrolling factor. Let's then try to set $k = 16$. This time, each thread computes its first index as

$$idx = blockIdx.x * blockDim.x * 16 + threadIdx.x * 16$$

and the fifteen following ones by summing 1, 2, ..., 15 to *idx*. The logic behind this implementation is unchanged with respect to the preceeding one.
The code for this algorithm can be found in
myProject/DIR___GPU_cp_dp/DIR___GPU_cp_dp_uroll, in the file *GPU_dp_uroll.cu*.
Its performances are reported in the following table. However, this time we didn't have an improvement, but instead a worsening.

Instance	Block dimension	Min	Max	Avg
custom_1	128	boh	bih	bah
custom_1	256	boh	bih	bah
custom_1	512	boh	bih	bah
custom_1	1024	boh	bih	bah

bla bla di nuovo le performance
This is probably due to the fact that, with such an unrolling factor, a single thread is tasked with too many memory operations that can't be parallelized completely, and so we lose parallelism.

4.4 SSD: Shared Memory

Another CUDA feature that allows programmers to achieve potential speedup in their applications is a program-managed, on-chip (so, per-Streaming Multiprocessor) cache. It is known as *Shared Memory* (SMEM), and it shares the same on-chip physical memory of the L1 cache. For this reason only, load/store operations performed by a thread to the

SMEM are much faster than those performed towards the global memory. However, to fully exploit its capabilities, one must know how data are organized and accessed in this memory.

When designing a kernel, a certain amount of Shared Memory can be declared, either statically or dynamically. During the execution, each *block* of the kernel has a certain amount of SMEM associated to it. This is an important point because the SMEM serves also as an intra-block communication mechanism. This means that threads of the same block can see and manage the same Shared Memory (the name *shared* it's not random), but have no access whatsoever to that of any other thread block. This has two consequences:

- Given that a single Streaming Multiprocessor can host multiple blocks at the same time, one should know how much SMEM can be given to a block in order not to affect the performances of the application. If a block requires a lot of SMEM, the hardware, that will associate that block to a SM, will also reserve enough Shared Memory for it, possibly preventing other blocks to be allocated on that same SM because of a lack of SMEM.

- When designing a Shared Memory based application, the programmer should reason about what the individual blocks of the grid will do, and how the threads will access to the SMEM.

The last point is particularly important because of the fact that the Shared Memory is organized in *Memory Banks*. Basically, the entire Shared Memory can be divided into 32 parts, or banks. When a warp issues a SMEM load/store request, this results more or less efficient depending on (once again) the number of memory transactions needed to satisfy such request. In general, the SMEM is organized in rows, and each row is composed of 32×4 bytes, where every 4 consecutive bytes belong to the same bank (let's call a group of 4 bytes belonging to the same bank *address*. Some architectures enlarge the bank width up to 8, althought this does not change the principles explained here). If a warp SMEM access is such that **at most** one address per bank is requested, then the memory request is said to be conflictless, and just one memory transaction is necessary to satisfy the overall request. This is the case every programmer should strive to achieve. Of course, a conflictless access to the Shared Memory is possible when the memory addresses requested are aligned (with respect to the shared memory rows) and coalesced, but in general, the important part consists in building a memory access pattern such that it is conflictless.

In our proposed implementation, each block requires $blockDim.x \times 2$ bytes of Shared Memory. In this way, every thread can move inside of it the two elements of the previous row necessary to compute the new one. If a thread is such that $idx < A[i]$, just $A[i-1]$ is copied in the SMEM. The actual address in which an element is copied depends on the *warp index* and the *thread index* of the current thread. The warp index is computed as

$$threadIdx.x/32$$

while the thread index is simply *threadIdx.x*.

Each warp uses 64 contiguous bytes of memory. The main operations performed are:

1. Load from the global memory the idx-th and (in case) the $(idx - A[i])$-th elements of the previous row, storing their values in the Shared Memory. Each warp uses the bytes that range from index $warp\_id \times 64$ included to $(warp\_id + 1) \times 64$ excluded, so that no data is overwritten during the lifetime of a kernel.

2. Then, the two values are read from the shared memory and put inside two local variables. If one element is not available, it is simply set to *false*.

3. Finally, the value for the *idx*-th entry is computed as $a \lor b$ and brought back to global memory.

Since each warp writes and reads a chunk of 64 contiguous bytes from the Shared Memory, with its starting index being a multiple of 64, the access is always coalesced and aligned, and can be satisfied with only one memory transaction.
The code for this algorithm can be found in
myProject/DIR___GPU_cp_dp/DIR___GPU_cp_dp_smem, in the file *GPU_dp_smem.cu*.
Its performances are reported in the following table.

Instance	Block dimension	Min	Max	Avg
custom_1	128	boh	bih	bah
custom_1	256	boh	bih	bah
custom_1	512	boh	bih	bah
custom_1	1024	boh	bih	bah

The overall result is not bad at all: for large instances, we even save some seconds.
bla bla fai vedere i risultati e qualcosa sulle performance, buttaci dentro anche i risultati di diversi blocchi.

4.5   SSD: Unrolling 8 + Shared Memory

So far in this chapter, the unrolling technique with unrolling factor $k = 8$ and the Shared Memory brought us the best speedup to the standard Dynamic Programming parallel (v2) approach. It comes natural to think that merging these two approaches can give us even better performaces. Let's see if that's true.
With the Shared Memory approach, each thread had to bring into the on-chip cache at most two bytes from the previous row. In total, a warp managed 64 bytes of the SMEM. In the unrolling approach with factor $k = 8$, each thread had to read up to 16 bytes from the previous row, in order to determine 8 new values. Basically, each warp managed 16 × 32 = 512 bytes.
Combining these approaches, we would have threads that bring 16 bytes of data from global memory to the Shared Memory, for a total of *blockDim.x* × 2 × 8 bytes of SMEM necessary for each block. Each warp still manages 512 bytes in total, so a minimum of four rows of SMEM are necessary to contain them. If these bytes are coalesced and aligned in the Shared Memory, exactly four rows need to be accessed by a warp to fulfill a memory request.
All that remains is to decide how to organize these data into the Shared Memory. Since a SMEM row is made of 32 × 4 = 128 bytes, and at most 512 bytes are read by a warp, we can organize the application in the following way:

1. Each warp of a block is given a contiguous set of 512 SMEM bytes, with the starting one being an even multiple of 512. This guarantees us that the accesses to the memory are coalesced and aligned.

2. A given thread has an index *idx = blockIdx.x * blockDim.x * 8 + threadIdx.x * 8*, that tells which is the first element of the new row that it must compute. The other ones are obtained by summing to *idx* the values 1, 2, ..., 7. Each of these indexes allow us to go to the previous row and read up to two values:

- The one of the same index;

- The one of the same index minus $A[i]$, if that index is a valid one.

These 16 values are read from the global memory and stored in the shared one, and each group of four is in a different row.

3. The bytes written in the SMEM are then brought in some local arrays (one for each thread), a and b, of lenght 8.

4. For $i = 0 \ldots 7$, we compute the new row's value of position $idx + i$ as

$$a[i] \vee b[i]$$

and store it into the global memory.

The code for this algorithm can be found in
myProject/DIR___GPU_cp_dp/DIR___GPU_cp_dp_uroll8_smem, in the file *GPU_dp_uroll8_smem.cu*.
Its performances are reported in the following table.

Instance	Block dimension	Min	Max	Avg
custom_1	128	boh	bih	bah
custom_1	256	boh	bih	bah
custom_1	512	boh	bih	bah
custom_1	1024	boh	bih	bah

The rusulting execution times show that, indeed, the Shared Memory - Unrolling 8 combo is effective, and outperforms the other ones. We could try doing the same with an unrolling factor of $k = 16$, but since that unrolling wasn't that effective in the first place, it probably wouldn't improve with the help of the Shared Memory. bla bla mostra i tempi di esecuzione

5 GPU (Parallel) algorithms: exhaustive search

The approach seen in the previous chapter is extremely good, but has one unsolvable problem: the memory limit. In order for the Dynamic Programming parallel approach to work, in fact, we need to allocate $(c + 1) \times 2$ bytes on the device global memory. If the instance is such that the capacity is greater than half the total size of the global memory, this approach cannot be used. However, one approach that didn't have any memory problem was the exhaustive one (since its space complexity was $\mathcal{O}(n)$). Althought intrinsically inefficient time-wise, it has no rivals space-wise. So, we could think of a parallel implementation of this approach.
Recursion in cuda is possible thanks to Dynamic Parallelism. In Dynamic Parallelism, a kernel's thread is able to launch a nested grid of thread blocks, and as such can be used to implement recursive functions. However, due to the limits of Dynamic Parallelism-based kernels, such as the maximum nested grid's depth, and the general complexity of designing a scalable application based on this feature, we won't rely on it to implement an exhaustive parallel search.

5.1 The abstract idea

What are we actually doing when performing an exhaustive search? Given the volumes and their number, n, we are taking, one by one, all the elements from the powerset of the volumes set, and for each of them we compute its value (given by the sum of the volumes considered). If we are lucky and smart enough, we can a priori exclude some of these sets because we recognize their unfeasibility. We did exactly this with the sequential exhaustive approach, where we cut the recursive exploration as soon as the capacity was exceeded.

An idea we came up with for a parallel approach is, in a very abstract way:

1. Partition the search space, given by 2^n elements, into subsets.

2. For each of them, find out if it *could* contain the optimal solution or not.

3. If it could, examine it exaustively. Otherwise, skip it.

To make up an algorithm from this idea, we have to answer two questions:

1. How should we partition the search space?

2. How can we find out if a subset **surely does NOT contain** the optimal solution?

5.2 The abstract algorithm

In order to answer in an effective way these questions, we have to describe what we actually want to do in our code. To do so, we have many different intuitions hat need to be exposed.

5.2.1 Intuition 1

Given that our search space is the powerset of the set of volumes, how can we explicitly represent and identify one of its elements (a tentative solution, a subset of the volumes set)? In the sequential implementation, the solutions were generated implicitly: when the bottom of the recursion was reached, the *accum* variable held all the volumes considered by the solution generated, and so, indirectly, also the ignored volumes. But in this case we want to explicitly say, in a tentative solution: "this volume is part of the sum; this other volume is not; this one more is part of the solution", etc.. How can we do that?

Since the presence or absence of a volume in a solution is a binary information, one bit will suffice.

When we deal with the binary representation of positive numbers, we have an interesting property: *the rightmost bits refer to lowest powers of two, while the leftmost ones refer to the highest powers of two.* We can define the order in which the solutions will be visited.

HAHAHA cattivo. 7 5 5 5 3 3

Index (base 2)	Index (base 10)	Subset sum	Value	Dist. 4	Dist. 8	Dist. 16	Dist. 32
0	000000	0+0+0+0+0+0	0	0	0	0	0
1	000001	0+0+0+0+0+3	3				
2	000010	0+0+0+0+3+0	3				
3	000011	0+0+0+0+3+3	6	6			
4	000100	0+0+0+5+0+0	5	5			
5	000101	0+0+0+5+0+3	8				
6	000110	0+0+0+5+3+0	8				
7	000111	0+0+0+5+3+3	11	11	11		
8	001000	0+0+5+0+0+0	5	5	5		
9	001001	0+0+5+0+0+3	8				
10	001010	0+0+5+0+3+0	8				
11	001011	0+0+5+0+3+3	11	11			
12	001100	0+0+5+5+0+0	10	10			
13	001101	0+0+5+5+0+3	13				
14	001110	0+0+5+5+3+0	13				
15	001111	0+0+5+5+3+3	16	16	16	16	
16	010000	0+5+0+0+0+0	5	5	5	5	
17	010001	0+5+0+0+0+3	8				
18	010010	0+5+0+0+3+0	8				
19	010011	0+5+0+0+3+3	11	11			
20	010100	0+5+0+5+0+0	10	10			
21	010101	0+5+0+5+0+3	13				
22	010110	0+5+0+5+3+0	13				
23	010111	0+5+0+5+3+3	16	16	16		
24	011000	0+5+5+0+0+0	10	10	10		
25	011001	0+5+5+0+0+3	13				
26	011010	0+5+5+0+3+0	13				
27	011011	0+5+5+0+3+3	16	16			
28	011100	0+5+5+5+0+0	15	15			
29	011101	0+5+5+5+0+3	18				
30	011110	0+5+5+5+3+0	18				
31	011111	0+5+5+5+3+3	21	21	21	21	21
32	100000	7+0+0+0+0+0	7	7	7	7	7
33	100001	7+0+0+0+0+3	10				
34	100010	7+0+0+0+3+0	10				
35	100011	7+0+0+0+3+3	13	13			
36	100100	7+0+0+5+0+0	12	12			
37	100101	7+0+0+5+0+3	15				
38	100110	7+0+0+5+3+0	15				
39	100111	7+0+0+5+3+3	18	18	18		
40	101000	7+0+5+0+0+0	12	12	12		
41	101001	7+0+5+0+0+3	15				
42	101010	7+0+5+0+3+0	15				
43	101011	7+0+5+0+3+3	18	18			
44	101100	7+0+5+5+0+0	17	17			
45	101101	7+0+5+5+0+3	20				
46	101110	7+0+5+5+3+0	20				
47	101111	7+0+5+5+3+3	23	23	23	23	

48	110000	7+5+0+0+0+0	12	12	12	12	
49	110001	7+5+0+0+0+3	15				
50	110010	7+5+0+0+3+0	15				
51	110011	7+5+0+0+3+3	18	18			
52	110100	7+5+0+5+0+0	17	17			
53	110101	7+5+0+5+0+3	20				
54	110110	7+5+0+5+3+0	20				
55	110111	7+5+0+5+3+3	23	23	23		
56	111000	7+5+5+0+0+0	17	17	17		
57	111001	7+5+5+0+0+3	20				
58	111010	7+5+5+0+3+0	20				
59	111011	7+5+5+0+3+3	23	23			
60	111100	7+5+5+5+0+0	22	22			
61	111101	7+5+5+5+0+3	25				
62	111110	7+5+5+5+3+0	25				
63	111111	7+5+5+5+3+3	28	28	28	28	28

Index base 2	Index base 10	Subset sum	Value	Distance 4
	1	2		
	1	2 Foo		
	1	2		
	1	2		
Bar	1	2		
	1	2		
	1	2		
	1	2		

References

[1] katema official. Gpu_computing_project. https://github.com/katema-official/GPU_Computing_Project.

[2] John Cheng, Max Grossman, and Ty McKercher. *Professional CUDA C Programming.* Wrox Press, 2014.

[3] Chad Parry. Reduction of subset-sum-optimization to subset-sum-decision. https://courses.cs.washington.edu/courses/csep521/05sp/lectures/sss.pdf.